

Metadata Internals

Copyright © 2001 by John Lam



About this talk

- This talk is about:
 - Metadata tables and tokens
 - System.Reflection and IMetaDataImport
 - Metadata on-disk format
 - Metadata scopes and reference resolution
- This talk is NOT about:
 - System.Reflection details
 - Dynamic binding
 - Reflection.Emit

Overview

- What is metadata?
- How do I use the reflection API?
- How do I use the unmanaged metadata API?
- What's different about these API's?
- How do I deal with metadata scopes?

What is metadata?

- Metadata: *n. data about data*
- Describes types
 - Member fields, methods, properties, events
- Describes signatures
 - Fields, properties, methods, delegates, local vars
- Describes types and *references*
- Describes miscellaneous entities
 - Files, modules, assemblies

Metadata in action

- Metadata by itself is useless
 - Like abstract base classes in C++ / Java / C#
 - Like IDL / TLB in COM
- CIL instructions reference metadata via **tokens**
 - First byte of token indicates type
 - Last three bytes of token are either row # (tables) or offsets (heaps)
 - Tokens are stored compressed in signatures (binary file viewers *not* very useful)

How to examine tokens

- Use ILDASM
 - Use /TOKENS option: displays token values
 - Use /BYTES option: displays CIL bytes
 - Use /OUT option: dump disassembly to disk file

```
ildasm /TOKENS /BYTES /OUT=HelloWorld.il HelloWorld.exe
```

How to examine metadata heaps

- Use MetaInfo
 - Fully documented sample metadata reader
 - In Tool Developers Guide\Samples\MetaInfo
- You must build it yourself!
 - See comments in metainfo.mak
 - Use short path names (GetShortPathName API)

```
metainfo /raw /heaps mscorlib.dll > dump.txt
```

Hello, World

```
class HelloWorld
{
    static void Main(string[] args)
    {
        Console.WriteLine( "Hello, world" );
    }
}
```



```
.method /*06000001*/ private hidebysig static
    void Main(string[] args) cil managed
{
    .entrypoint
    .maxstack 8
    IL_0000:  /* 72      | (70)000001 */ ldstr    "Hello, world"      /* 70000001 */
    IL_0005:  /* 28      | (0A)00000E */ call    void [mscorlib /* 23000001 */
                                     System.Console /* 0100000F */
                                     ::WriteLine(string) /* 0A00000E */
    IL_000a:  /* 2A      |          */ ret
} // end of method HelloWorld::Main
```


About HelloWorld.exe

- Five tokens (and types) in listing
 - MethodDef: (0x06000001)
 - User string: (0x70000001)
 - AssemblyRef: (0x23000001)
 - TypeRef: (0x0100000F)
 - MemberRef: (0x0A00000E)
- Only two are directly referenced by CIL in this example:
 - String (by **ldstr** instruction)
 - MemberRef (by **call** instruction)

What's a string token?

- String tokens refer to strings
 - Offset into user string (#US) heap
- Points to UTF-8 string
 - Length prefixed (PackedLen) (includes null)
 - Terminated by single-byte null
- Used only by **ldstr** instruction

Definition vs. reference tokens

- Two general categories of metadata tokens
 - Definition tokens
 - Reference tokens
- Definition tokens define the entity
 - e.g. TypeDef tokens describe definition of Types
- Reference tokens define references to the entity
 - e.g. TypeRef token describes how to find TypeDef

A MemberRef is logically ...

- MemberRefs reference "type members"
 - Fields, Methods
 - Colloquially referred to as FieldRef, MethodRef
- What information does a MemberRef need?
 - Member's signature (think method overloading)
 - Type that member is found in
 - Assembly that type is found in

A MemberRef is physically ...

- Dereferencing MemberRef token yields table row
- MemberRef row contains
 - Parent (typically a TypeRef token)*
 - Name (offset in #Strings heap)
 - Signature (offset in #Blob heap)
- * simplified for purposes of discussion
- Documented in ECMA Partition II spec

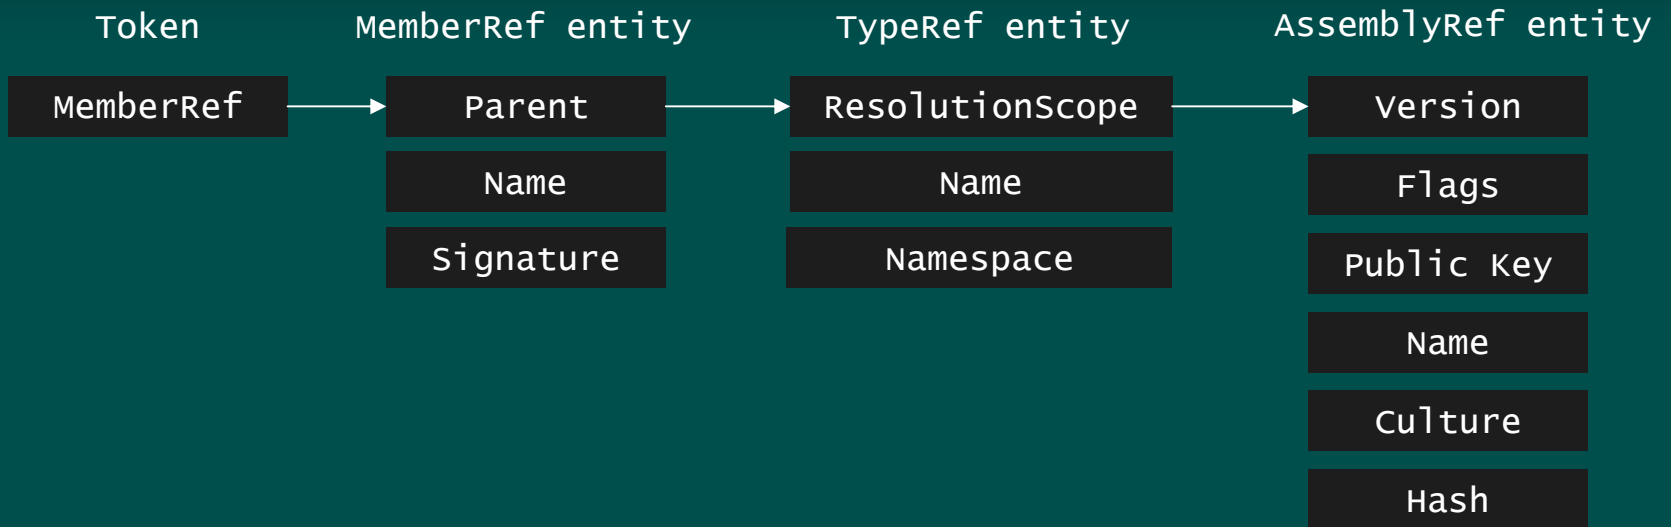
TypeRef tokens

- Dereferencing TypeRef token yields TypeRef row
- TypeRef row contains
 - ResolutionScope (typically an AssemblyRef token)*
 - Name (string)
 - Namespace (string)
- * simplified for purposes of discussion

AssemblyRef tokens

- AssemblyRef row contains
 - Major, Minor, Build, Revision #'s (ushort)
 - Flags (uint)
 - Public key (blob)
 - Name (string)
 - Culture (string)
 - Hash value (blob)

MemberRef



Metadata scopes

- Metadata definition tokens scoped by module
 - TypeDef, MethodDef, AssemblyDef ...
- Metadata reference tokens scoped by app domain
 - TypeRef, MemberRef, AssemblyRef, ModuleRef
- CLR loads into app domain:
 - Statically referenced assemblies at start-up
 - Dynamically referenced assemblies as needed

Metadata API's

- Two major metadata API's available
 - Managed: System.Reflection
 - Unmanaged: IMetaDataImport and friends
- Both are limited in different ways
 - System.Reflection does not permit access to CIL
 - System.Reflection does not reveal token values
 - IMetaDataImport cannot resolve *Ref tokens

System.Reflection

- API for inspecting metadata
 - Provides a *logical* view of metadata
 - Physical details (e.g. token values, CIL) obscured
- API for emitting metadata (and CIL)
 - System.Reflection.Emit

Loaded assemblies

- AppDomain object returns array of assemblies

```
private static void EnumerateLoadedAssemblies()
{
    Assembly[] assemblies = AppDomain.CurrentDomain.GetAssemblies();
    foreach( Assembly assembly in assemblies )
        Console.WriteLine( assembly.FullName );
}
```

Referenced assemblies

- Assembly object returns array of AssemblyNames
- AssemblyName == AssemblyRef
 - Version, Name, Culture, Public Key, Hash

```
private static void EnumerateReferencedAssemblies( Assembly assembly )
{
    AssemblyName[] referencedAssemblies = assembly.GetReferencedAssemblies();
    foreach( AssemblyName assemblyName in referencedAssemblies )
        Console.WriteLine( "-- {0}", assemblyName.FullName );
}
```

Types in an assembly

- Assembly object returns array of Type objects
 - Struct, Enum, Class, Interface

```
private static void EnumerateTypes( Assembly assembly )
{
    Type[] types = assembly.GetTypes();
    foreach( Type type in types )
        Console.WriteLine( "== {0}", type.FullName );
}
```

Methods in a type

- Type object exposes array of MethodInfo objects
- MethodInfo represents method
 - Signature, return value, visibility ...

```
private static void EnumerateMethods( Type type )  
{  
    MethodInfo[] methods = type.GetMethods();  
    foreach( MethodInfo method in methods )  
        Console.WriteLine( "**** {0}", method.Name );  
}
```

Method's CIL

- System.Reflection does not permit access to CIL
 - API would need to understand / expose tokens
 - API would be significantly more complex
- Need to use a different API

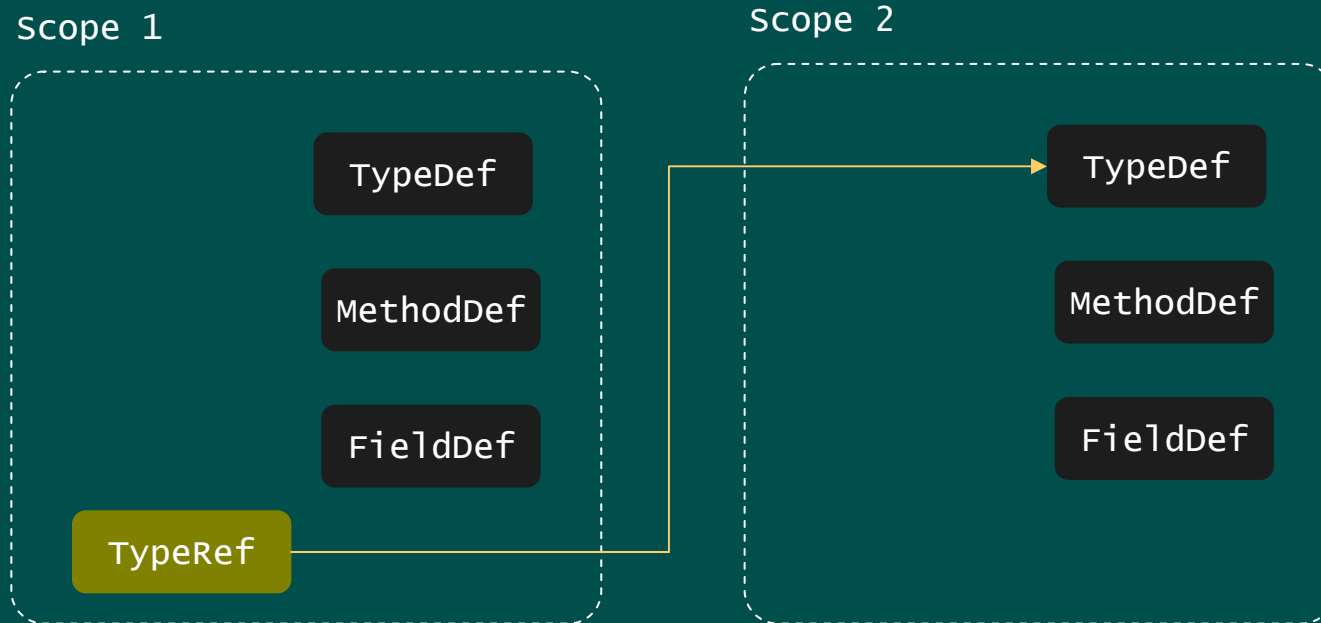
Unmanaged metadata API

- COM-based API
 - CLSID_CorMetaDataDispenser implementation
 - Initial interface is IMetaDataDispenserEx

```
hr = CoCreateInstance( CLSID_CorMetaDataDispenser,  
                       0,  
                       CLSCTX_INPROC_SERVER,  
                       IID_IMetaDataDispenserEx,  
                       reinterpret_cast< void** >( &_spMDD )  
                       );
```

Metadata scopes

- Dispensers enable access to *scopes*
 - Scopes define boundary for metadata definitions
 - Scopes are usually found in physical files



Opening a scope

- Open scope using OpenScope() method
 - Returns IMetaDataImport interface
 - Each module has its own scope

```
SIMetaDataImport OpenScope( SIMetaDataDispenserEx dispenser, string path )
{
    USES_CONVERSION;
    SIMetaDataImport result;

    dispenser->OpenScope( A2W( path.c_str() ),
                          0,                                     // read
                          IID_IMetaDataImport,
                          reinterpret_cast< IUnknown** >( &result )
    );

    return result;
}
```

Enumerators

- Enumerators iterate over items in a scope
 - Identified by HCORENUM handle
 - Types, Methods, Fields, Parameters ...
- Usage: acquire enumerator, use, close

```
void Enumerate( SIMetaDataImport metadata )
{
    HRESULT    hr;
    HCORENUM    hEnum = 0;
    hr = metadata->EnumXXX( &hEnum, // [in/out] enumerator handle
                           ...
                           );
    if( SUCCEEDED( hr ) )
    {
        // Use
        metadata->CloseEnum( hEnum );
    }
}
```

Enumerating types

- Use EnumTypeDefs() method to retrieve tokens

```
void EnumTypes( SIMetaDataImport metadata )
{
    HRESULT    hr;
    DWORD      count;
    HCORENUM    hEnumTypeDef = 0; // NOTE that this is an in/out param and MBZ!
    mdTypeDef  typeDefs[ MAX_TYPEDEFES_PER_MODULE ];

    hr = metadata->EnumTypeDefs( &hEnumTypeDef,           // [in/out] enumerator handle
                                typeDefs,                 // [out] array of mdTypeDefs
                                NumItems( typeDefs ),     // [in] size of array
                                &count,                   // [out] number of items
                                );

    if( SUCCEEDED( hr ) )
    {
        for( unsigned int i = 0; i < count; i++ )
            cout << GetTypeNames( typeDefs[ i ], metadata ).c_str() << endl;
        metadata->CloseEnum( hEnumTypeDef );
    }
}
```

Enumerating types

- Use GetTypeDefProps() to retrieve token data

```
string GetTypeNames( mdTypeDef typeDef, SIMetaDataImport import )
{
    USES_CONVERSION;
    HRESULT hr;
    ULONG    sizeName;
    wchar_t  wszName[ MAX_SYMBOL_LENGTH ];

    hr = import->GetTypeDefProps( typeDef,                // [in] typeDef token
                                   wszName,                // [out] buffer for name
                                   NumItems( wszName ),     // [in] size of name buffer
                                   &sizeName,              // [out] size of name
                                   0,                      // [out] CorTypeAttrs
                                   0,                      // [out] type it extends
                                   );

    _ASSERT( SUCCEEDED( hr ) );
    _ASSERT( sizeName < NumItems( wszName ) );

    return W2A( wszName );
}
```

Enumerating methods

- Use EnumMethods() to retrieve tokens

```
void EnumMethods( SIMetaDataImport metadata, mdTypeDef type )
{
    HRESULT      hr;
    DWORD        count;
    HCOENUM      hEnumMethods = 0;
    mdMethodDef methodDefs[ MAX_METHODDEFS_PER_TYPE ];

    hr = metadata->EnumMethods( &hEnumMethods,           // [in/out] enumerator handle
                               type,                     // [in] typeDef
                               methodDefs,               // [out] mdMethodDefs array
                               NumItems( methodDefs ),   // [in] size of array
                               &count                  // [out] number of items
                               );

    if( SUCCEEDED( hr ) )
    {
        for( unsigned int i = 0; i < count; i++ )
            cout << GetMethodName( methodDefs[ i ], metadata ).c_str() << endl;
        metadata->CloseEnum( hEnumMethods );
    }
}
```

Enumerating methods

- Use `GetMethodProps()` to retrieve token data

```
string GetMethodName( mdMethodDef methodDef, SIMetaDataImport import )
{
    USES_CONVERSION;
    HRESULT hr;
    ULONG    sizeName;
    wchar_t  wszName[ MAX_SYMBOL_LENGTH ];
    hr = import->GetMethodProps( methodDef,                // [in] methodDef token
                                0,                        // [out] typeDef token
                                wszName,                  // [out] buffer for name
                                NumItems( wszName ),       // [in] size of name buffer
                                &sizeName,                // [out] size of name
                                0,                        // [out] CorMethodAttrs
                                0,                        // [out] ptr to sig blob
                                0,                        // [out] size of sig blob
                                0,                        // [out] RVA to code
                                0                         // [out] implementation flags
                                );
    _ASSERT( SUCCEEDED( hr ) );
    _ASSERT( sizeName < NumItems( wszName ) );

    return W2A( wszName );
}
```


RVA to offset conversion

- Start of method's CIL determined by RVA
 - Relative Virtual Address
- Convert to offset into PE file
 - ImageRvaToVa() in dbghelp.lib

```
pHeader = (unsigned __int8*) ImageRvaToVa( g_File->getNTHeaders(),  
                                           g_File->getBase(),  
                                           rva,  
                                           0  
                                           );
```

Examining CIL

- CIL stream contains two sections
 - Header
 - CIL instructions
- Header depends on size of CIL stream
 - First byte determines type of header
 - CorILMethod_TinyFormat
 - CorILMethod_FatFormat

Examining CIL

□ Convert to byte pointer

```
unsigned __int8* GetMethodPointer( ULONG rva, ULONG* codeSize )
{
    unsigned __int8*      pHeader;
    IMAGE_COR_ILMETHOD_FAT* pFatMethodHeader;
    pHeader = (unsigned __int8*) ImageRvaToVa( g_File->getNTHeaders(),
                                              g_File->getBase(), rva, 0 );
    switch( *pHeader & CorILMethod_FormatMask )
    {
    case CorILMethod_TinyFormat:
        *codeSize = ((TinyHeader*)pHeader)->size;
        return ++pHeader;
    case CorILMethod_FatFormat:
        pFatMethodHeader = reinterpret_cast< IMAGE_COR_ILMETHOD_FAT* >( pHeader );
        *codeSize = pFatMethodHeader->CodeSize;
        return reinterpret_cast< unsigned __int8* >( pFatMethodHeader )
            + pFatMethodHeader->Size * sizeof( DWORD );
    }
}
```

Examining CIL

□ Enumerate over byte pointer

```
void DumpCIL( SIMetaDataImport metadata, mdMethodDef methodDef )
{
    HRESULT hr;
    ULONG rva;
    hr = metadata->GetMethodProps( methodDef, 0, 0, 0, 0, 0, 0, 0,
                                   &rva, 0 );
    if( SUCCEEDED( hr ) )
    {
        ULONG codeSize;
        unsigned __int8* pMethodBytes = GetMethodPointer( rva, &codeSize );

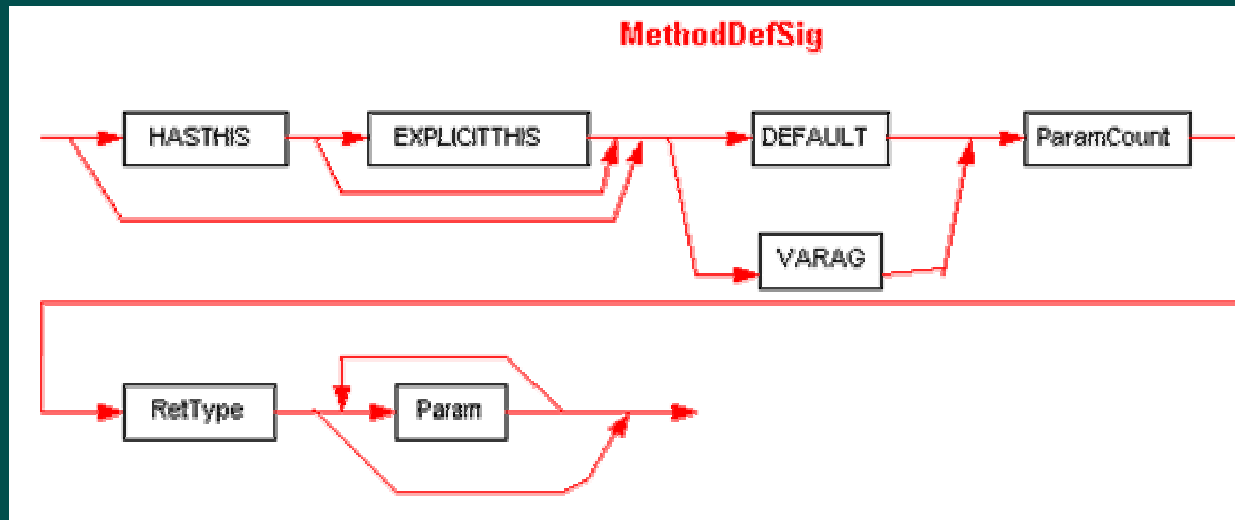
        for( ULONG i = 0; i < codeSize; i++ )
        {
            char szOutput[ 254 ];
            wsprintf( szOutput, "%2.2x", *pMethodBytes++ );
            cout << szOutput << " ";
        }
        cout << endl;
    }
}
```

Signatures and blobs

- Signatures describe many elements
 - Methods, fields, properties, local variables ...
- Signatures are variable-length data structures
 - Stored in metadata blob heap (#Blob)
- Method signatures describe
 - Parameters, calling convention, return values ...

Cracking method signatures

- Method signatures expose
 - Calling convention
 - Parameter list
 - Return values



Cracking method signatures

```
MethodDefSig::MethodDefSig( PCCOR_SIGNATURE sig, ULONG sigSize )
{
    PCCOR_SIGNATURE sigPos = sig;
    sigPos += CorSigUncompressData( sigPos, &_methodSig );    // read flags
    sigPos += CorSigUncompressData( sigPos, &_paramCount );   // read count

    _retvalSig = new RetvalSig( sigPos );    // crack return value sig
    sigPos += _retvalSig->TypeSize();        // advance pointer

    for( int i = 0; i < _paramCount; i++ )
    {
        ParamSig* paramSig = new ParamSig( sigPos ); // crack parameter sig
        sigPos += paramSig->TypeSize();                // advance pointer
        _paramSigs.push_back( paramSig );              // add to param sig array
    }

    _ASSERT( sigPos == sig + sigSize );    // sanity check
}
```

Uncompressing tokens

- Tokens are stored using compressed format
- Bits 0, 1 indicate token type
 - TypeDef (00), TypeRef (01), TypeSpec (10)
- Bits 2-31 stored as compressed integer
- Type info for all other token types discarded
 - Must discern token type from context

```
private static uint CorSigUncompressToken( BinaryReader reader )  
{  
    uint token = CorSigUncompressData( reader );  
    return (token >> 2) | (uint)CorEncodeToken[ token & 0x03 ];  
}
```


Uncompressing tokens

- Tokens are stored in a compressed format

```
private static uint CorSigUncompressData( BinaryReader reader ) {  
    uint result;  
    byte first = reader.ReadByte();  
    if( (first & 0x80) == 0x00 ) {           // 0xxx xxxx == 1 byte case  
        return (uint)first;  
    }  
    else if( (first & 0xc0) == 0x80 ) {      // 10xx xxxx == 2 byte case  
        result = (uint)((first & 0x3F) << 8);  
        result += reader.ReadByte();  
    }  
    else if( (first & 0xe0) == 0xc0 ) {      // 110x xxxx == 4 byte case  
        result = (uint)((first & 0x1F) << 24);  
        result += (uint)(reader.ReadByte() << 16);  
        result += (uint)(reader.ReadByte() << 8);  
        result += (uint)reader.ReadByte();  
    }  
    else  
        throw new Exception( "Illegal compressed token: first byte = {0}", first );  
    return result;  
}
```

Resolving Ref tokens

- Ref tokens resolve to Def tokens in foreign scope
- Foreign scope is identified in metadata
 - AssemblyRef provides information
- Need to *resolve* foreign scope
 - Find assembly file on disk
 - Download on-demand assemblies
- Need to use Fusion API
 - Not doc'd in current form

Resolving MemberRef



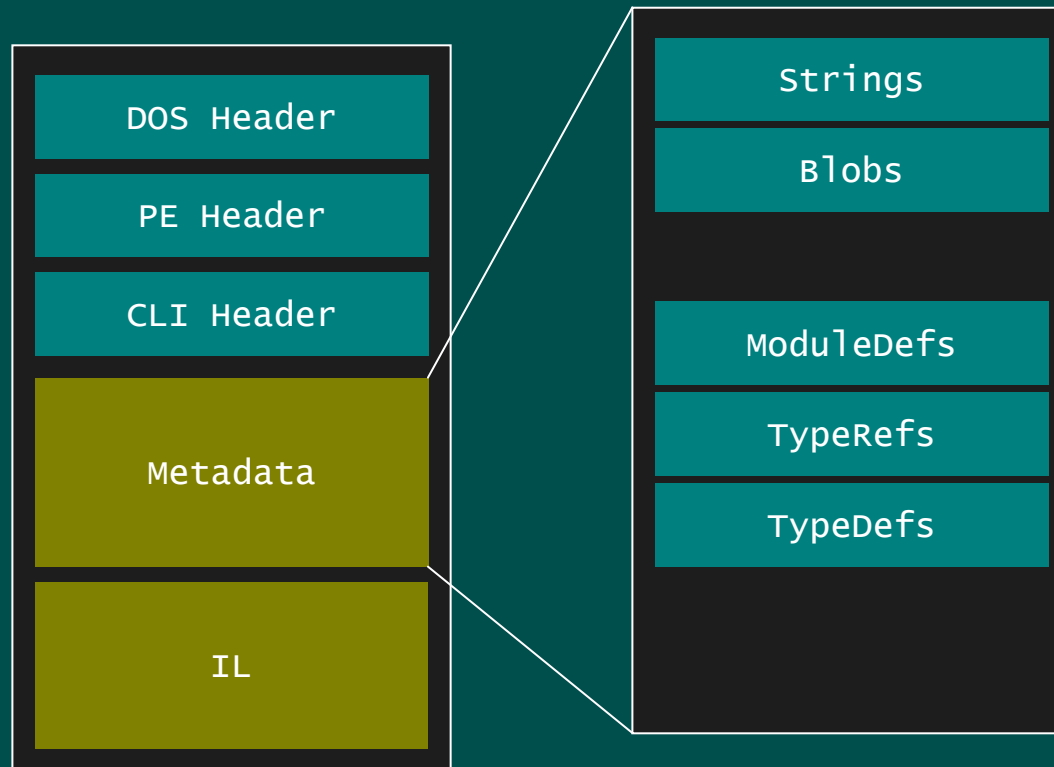
A better metadata API

- Combines benefits of managed, unmanaged API
 - Access to CIL
 - Access to tokens
 - System.Reflection ease of use
- Need to roll your own
 - Must understand metadata on-disk format

How is metadata stored?

- Metadata always stored in an assembly
 - Assemblies are usually single files
 - Single file is portable executable (PE) file
 - Custom CLI header contains metadata offset/size
- Metadata stored in tables / heaps
 - Tables indexed by row number
 - Heaps indexed by offset

How is metadata stored



Metadata on-disk format

- Documented in ECMA Partition II spec
 - CLI header format
 - Metadata directory format
 - Metadata table structures
- Metadata on-disk format is highly compressed
 - Optimized for size over speed

Resolving by name

- Resolving reference tokens done *by name*
 - Token elements contain a String token
- Resolving each can lead to different types
 - MemberRef -> MethodDef, PropertyDef, FieldDef, EventDef
 - TypeRef -> TypeDef
 - AssemblyRef -> AssemblyDef
 - ModuleRef -> ModuleDef

Resolving *Ref tokens

- Resolving *Ref token should yield corresponding *Def token
- Cannot do from System.Reflection
 - No access to CIL anyways ...
- Cannot do from IMetaDataImport
 - No access to Fusion API
- Solution requires using BOTH API's

Resolving MemberRefs



Resolving AssemblyRefs

- Problem:
 - Given an AssemblyRef
 - Return me the path to the assembly
- Solution:
 - Assembly.Load() solves this in managed API
 - Unmanaged API has no equivalent

Loading referenced assemblies

- `Assembly.Load()` used to locate assembly
 - `AssemblyName` is parameter
 - Uses internal `Load` method to do dirty work
 - Probably calls underlying Fusion API for resolution

```
public Assembly Load( AssemblyName assemblyRef );
```

References

- ECMA Partition II (Metadata) Specification
 - <http://msdn.microsoft.com/net/ecma>
 - Get PDF version!
- Unmanaged Metadata API
 - Tool Developer's Guide\docs directory